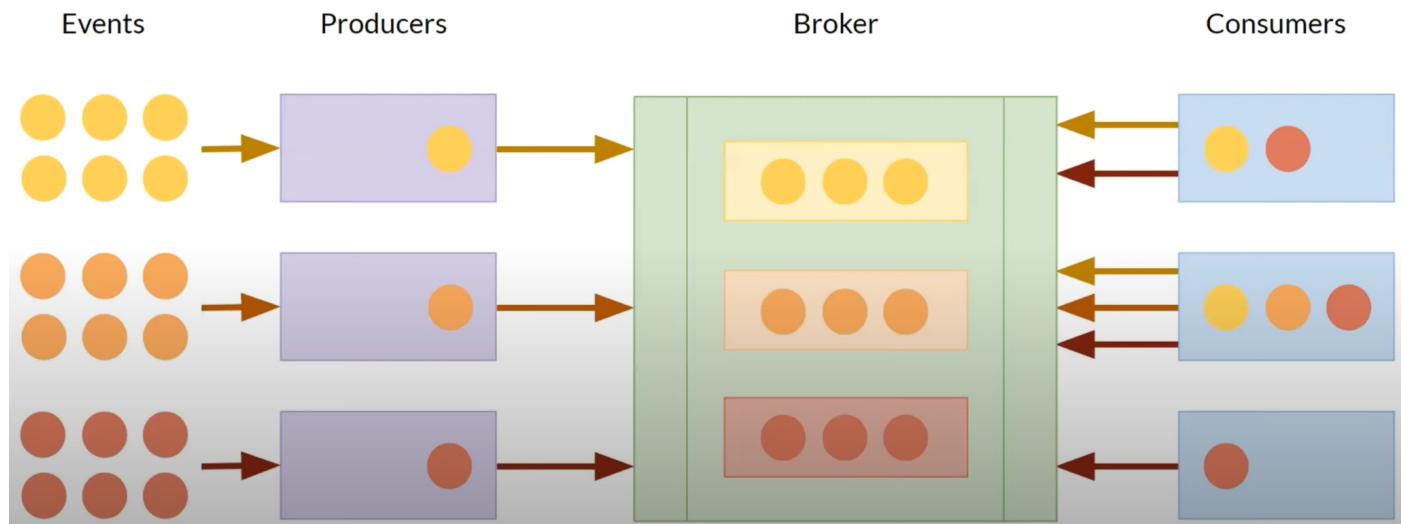


# Kafka

<https://www.youtube.com/watch?v=-AZ0i3kP9Js>

<https://www.youtube.com/watch?v=BtmYjTO1Epl>

Kafka is an open-source **message broker** project developed by the Apache Software Foundation written in Scala and is a distributed **publish-subscribe messaging system**. Kafka's design pattern is mainly based on the **transactional logs** design.



What can you do with Kafka?

- In order to transmit data between two systems, we can build a real-time stream of data pipelines with it.
- Also, we can build a real-time streaming platform with Kafka, that can actually react to the data.

Feature	Description
High Throughput	Support for millions of messages with modest hardware
Scalability	Highly scalable distributed systems with no downtime
Replication	Messages are replicated across the cluster to provide support for multiple subscribers and balances the consumers in case of failures
Durability	Provides support for persistence of message to disk
Stream Processing	Used with real-time streaming applications like Apache Spark & Storm

Data Loss	Kafka with proper configurations can ensure zero data loss
-----------	--

List the various components in Kafka:

- Topic – a stream of messages belonging to the same type
- Producer – that can publish messages to a topic
- Brokers – a set of servers where the published messages are stored
- Consumer – that subscribes to various topics and pulls data from the brokers.

## Zookeeper

Kafka uses **Zookeeper** to store offsets of messages consumed for a specific topic and partition by a specific Consumer Group. (Message offset role: messages contained in the partitions are assigned a unique ID number that is called the offset. The role of the offset is to uniquely identify every message within the partition)

It is not possible to bypass Zookeeper and connect directly to the Kafka server. If, for some reason, ZooKeeper is down, you cannot service any client request.

Apache ZooKeeper is an open-source server for highly reliable distributed coordination of cloud applications. It is a project of the Apache Software Foundation.

ZooKeeper is essentially a service for distributed systems offering a hierarchical **key-value store**, which is used to provide a distributed configuration service, synchronization service, and naming registry for large distributed systems. ZooKeeper was a sub-project of Hadoop but is now a top-level Apache project in its own right.

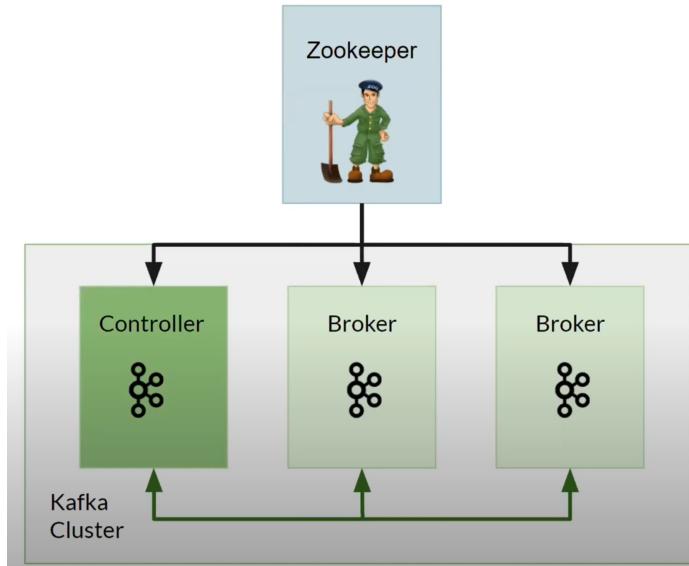
Typical use cases for ZooKeeper are:

- Naming service
- Configuration management
- Data Synchronization
- Leader election
- Message queue
- Notification system

Kafka message contains:

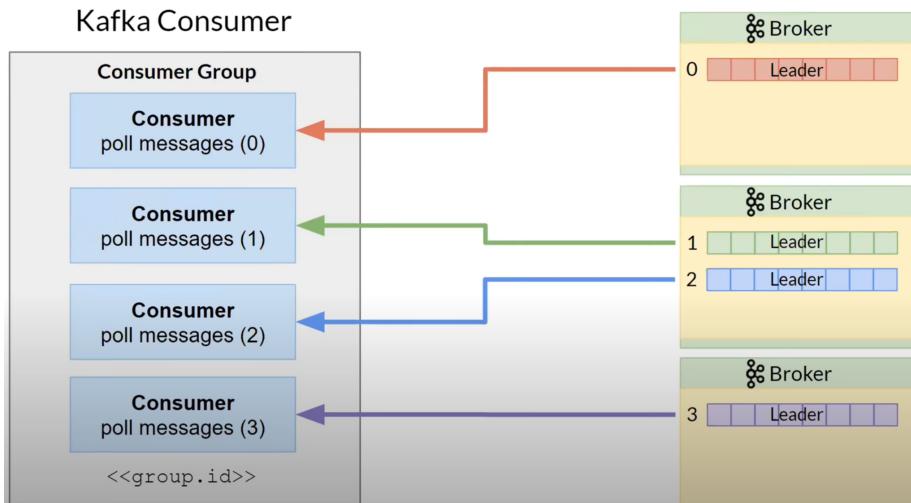
- key (optional) - used for message distribution across partitions
- value - byte array
- timestamp - set either before sending by client or after receiving by broker
- headers - key-value pairs/attributes

## Kafka Cluster



We may want to process messages with multiple consumers. We can organize concurrent processing of a topic using multiple partitions, and each partition is attached to its consumers using a group mechanism, as shown below.

- Single consumer can process more than 1 partition in a topic when **NoPartitions > NoConsumers**.
- Some consumers stay idle (have no job to do) when **NoPartitions < NoConsumers**

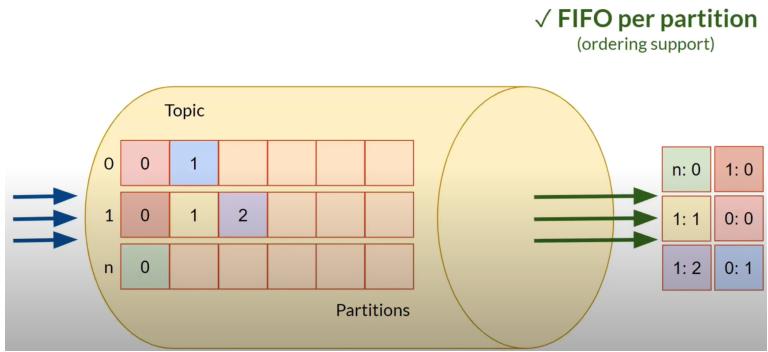


In order to overcome the challenges of collecting the large volume of data, and analyzing the collected data we need a messaging system. Hence Apache Kafka came into the story. Its benefits are:

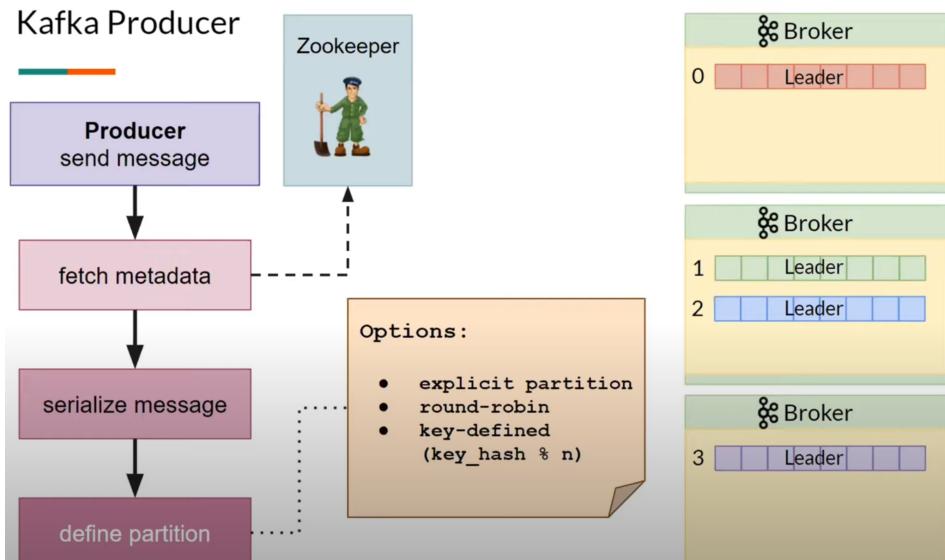
- It is possible to track web activities just by storing/sending the events for real-time processes.
- Through this, we can Alert as well as report the operational metrics.
- Also, we can transform data into the standard format.

- Moreover, it allows continuous processing of streaming data to the topics.

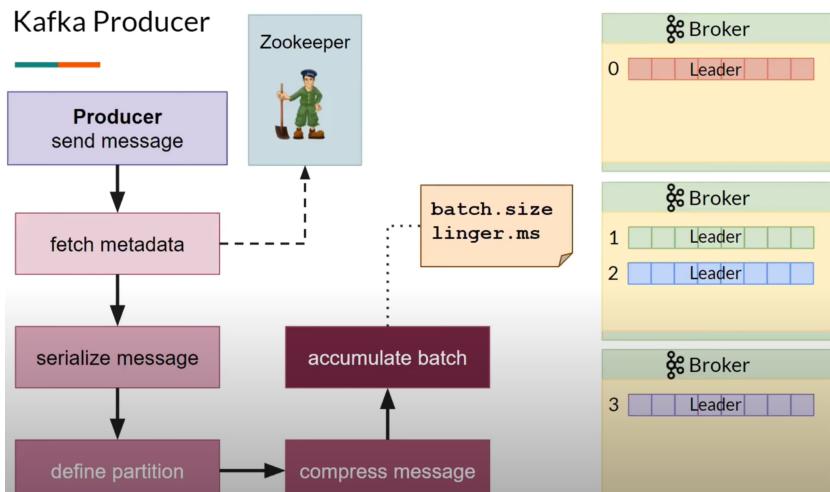
# Topics and partitions



Choosing partition to send the message to:



Sending batch to broker:



## Topic VS Queue

Is Kafka a Topic or a Queue?

<https://abhishek1987.medium.com/kafka-is-it-a-topic-or-a-queue-30c85386af6>

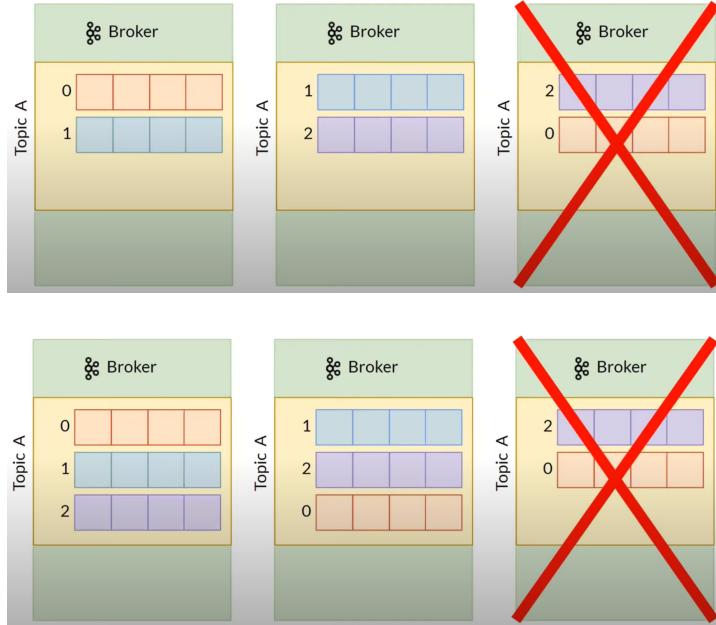
Kafka is both a Topic and a Queue. Queue based systems are typically designed in a way that there are multiple consumers processing data from a queue and the work gets distributed such that **each consumer gets a different set of items to process**. Hence there is no overlap, allowing the workload to be shared and enables horizontally scalable architectures.

A Kafka topic is sub-divided into units called partitions for fault tolerance and scalability. **Consumer Groups allow Kafka to behave like a Queue**, since each consumer instance in a group processes data from a non-overlapping set of partitions (within a Kafka topic).

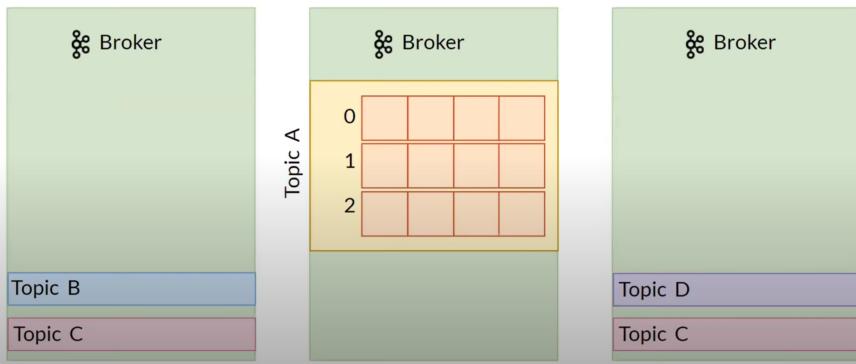
# Replication

Replicas are essentially a list of nodes that **replicate the log for a particular partition** irrespective of whether they play the role of the Leader. On the other hand, ISR stands for In-Sync Replicas. It is essentially a set of message replicas that are synced to the leaders.

Replication ensures that published messages are not lost and can be consumed in the event of any machine error, program error or frequent software upgrades.

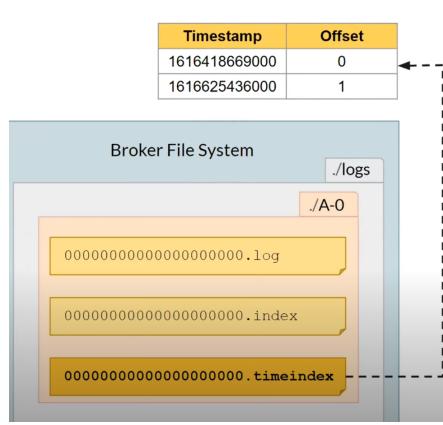
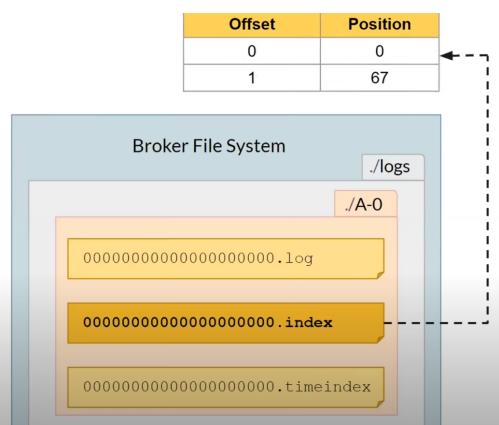
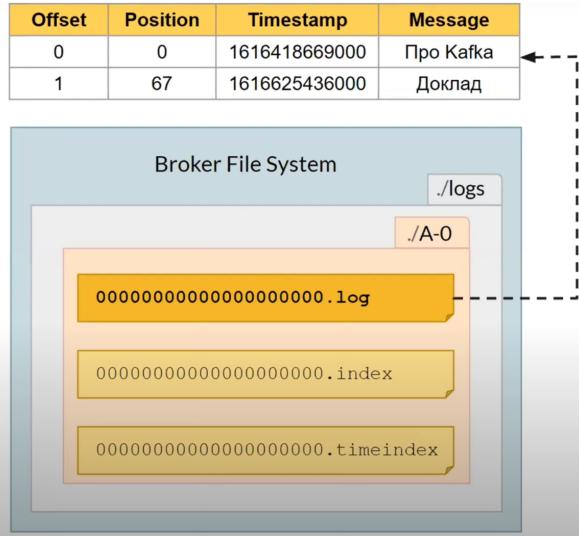
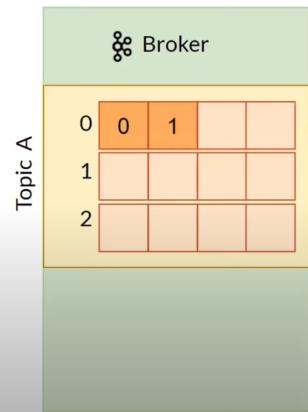


Problem of imbalance partition placement be like:

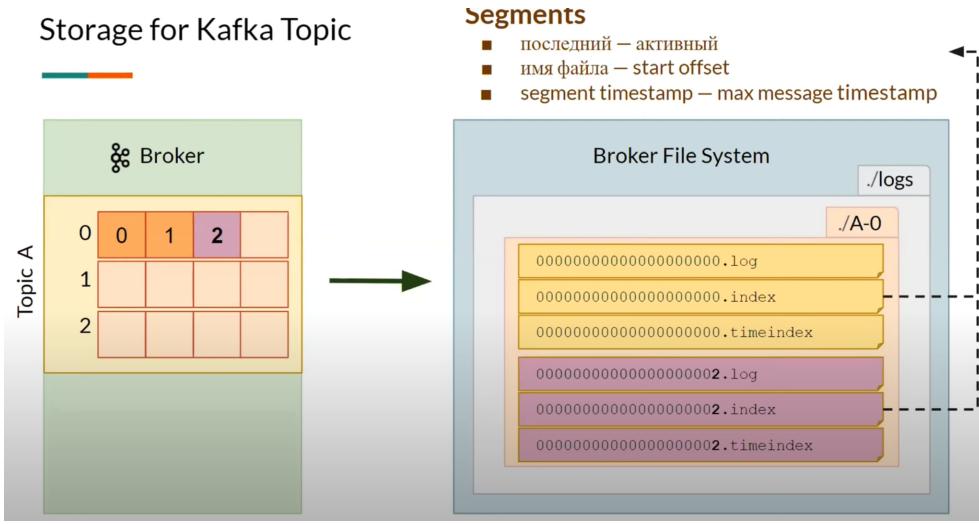


Data is stored in log files:

## Storage for Kafka Topic



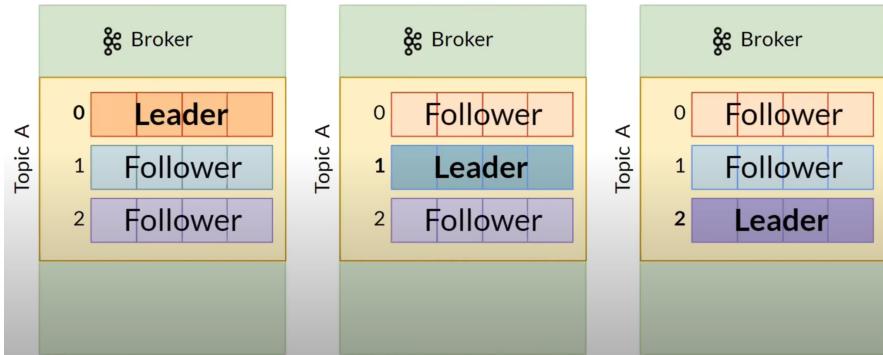
Every .log file size is limited:



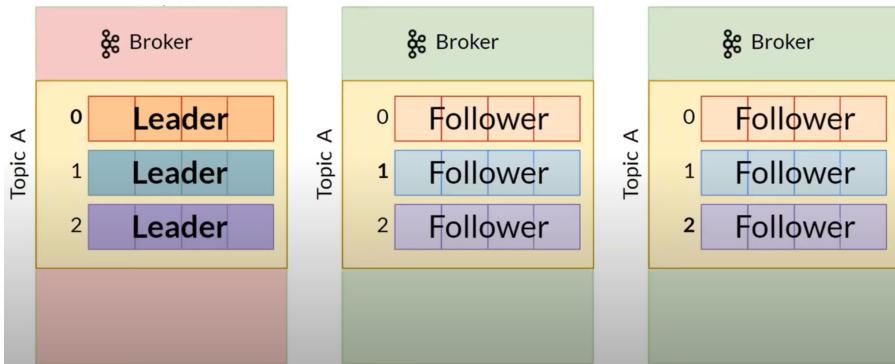
## Leader-Follower

Every partition in Kafka has one server which plays the role of a Leader, and none or more servers that act as Followers. **The Leader performs the task of all read and write requests** for the partition, while the role of the Followers is to passively replicate the leader. In the event of the Leader failing, one of the Followers will take on the role of the Leader. This ensures load balancing of the server.

replication-factor: 3



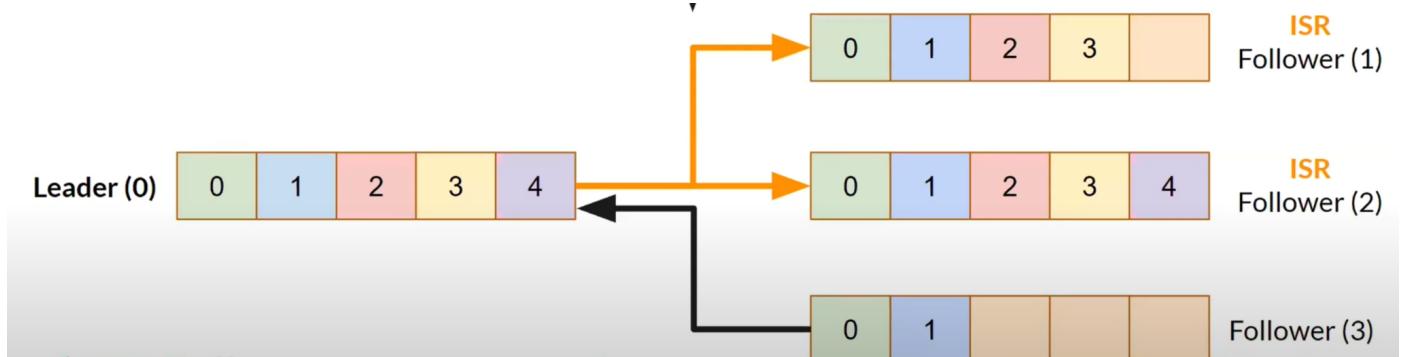
Imbalance leader election problem (Broker #1 is under heavy load)



## Data loss

A data loss can happen during leader reelection even if you have replication factor > 1, because generally speaking replicas are out-of-sync with the leader. In order to solve this problem, kafka supports In-Sync-Replicas (ISR).

Leader makes sure all the ISR replicas are in-sync before sending 'ack'-response to the producer. ISRs are good candidates for a leader reelection.



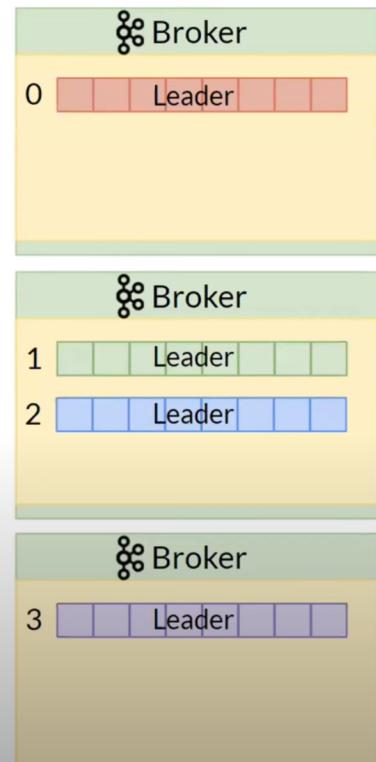
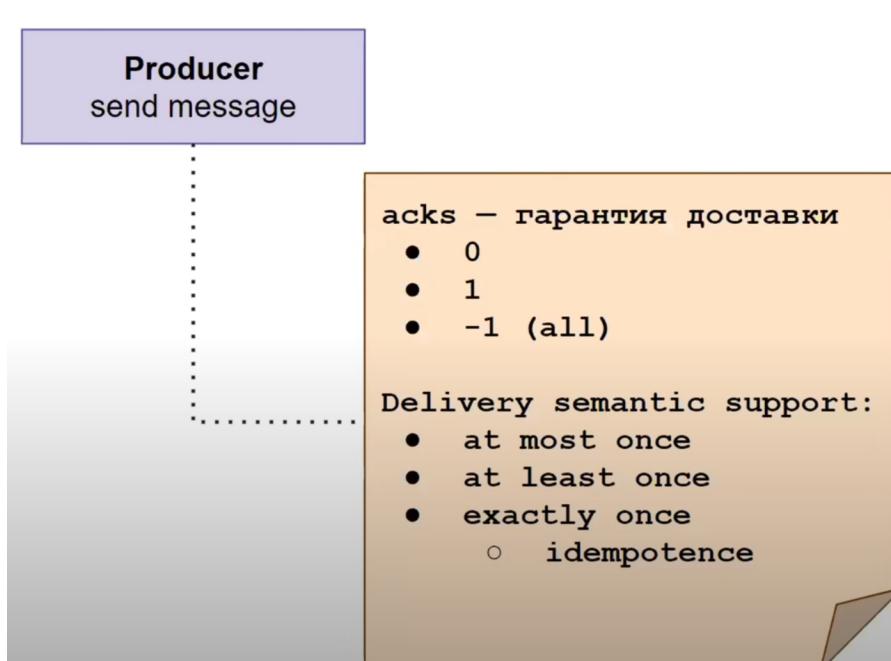
## Delivery semantics

ack = 0 - no confirmation from broker required

ack = 1 - leader has processed the message

ack = -1 (all) - leader and all ISR have processed the message

## Kafka Producer



What guarantees of delivery/message processing do we have?

## Kafka Consumer Offset commit

---

### Types of commits

- Auto commit
  - at most once (miss messages)
- Manual commit
  - at least once (duplicate messages)

**At most once** scenario (auto commit):

Consumer pulls message -> auto commit happens (update an offset for a group in \_\_consumer\_offset topic) -> consumer fails -> consumer handles the message.

The offset is updated but the message has not been processed. **The message is lost.**

**At least once** scenario (manual commit):

Consumer pulls message -> consumer handles the message -> consumer fails -> consumer manually commits.

After the consumer restarts the same message would be pulled and processed all over again. **Duplicate processing.**

**Exactly once** scenario:

Implement a mechanism of storing offset in an external storage, do not rely on \_\_consumer\_offset data.

Example:

Event X = money transfer from User A to User B. We want to handle this event exactly once, obviously. Suppose we use RDBMS that supports transactions. In that DB, we store info about user accounts as well as the kafka offset value for the most recent operation processed. After reading Event X, we perform a money transfer and update the offset value within a single transaction.

## Performance

### Kafka performance

---

#### Why is it so fast?

- ✓ Scalable architecture
- ✓ Sequential write and read
- ✓ No random read
- ✓ Zero-copy
- ✓ Huge amount of settings for different cases

## Starting kafka

Typical docker-compose.yml running zookeeper + kafka:

```
version: '2'
services:
  zookeeper:
    container_name: zookeeper
    networks:
      - kafka_network
    ...
  kafka:
    container_name: kafka
    networks:
      - kafka_network
    ports:
      - 29092:29092
    environment:
      KAFKA_LISTENERS: EXTERNAL_SAME_HOST://:29092,INTERNAL://:9092
      KAFKA_ADVERTISED_LISTENERS:
        INTERNAL://kafka:9092,EXTERNAL_SAME_HOST://localhost:29092
        KAFKA_LISTENER_SECURITY_PROTOCOL_MAP:
          INTERNAL:PLAINTEXT,EXTERNAL_SAME_HOST:PLAINTEXT
        KAFKA_INTER_BROKER_LISTENER_NAME: INTERNAL
      ...
  networks:
    kafka_network:
      name: kafka_docker_example_net
```

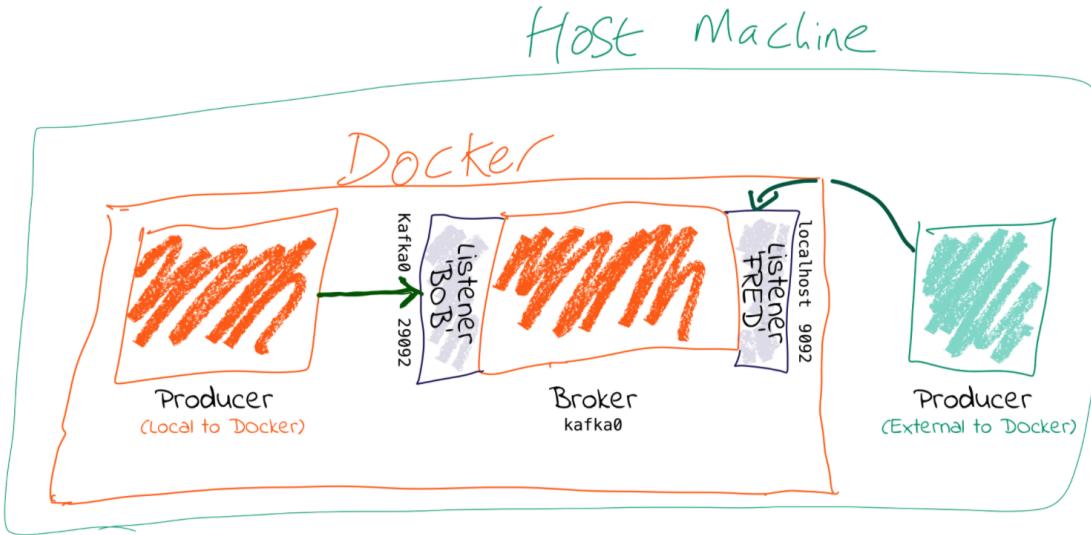
## Kafka Listeners explained

<https://rmoff.net/2018/08/02/kafka-listeners-explained/>

Example config:

```
KAFKA_LISTENERS: LISTENER_BOB://kafka0:29092,LISTENER_FRED://localhost:9092
KAFKA_ADVERTISED_LISTENERS: LISTENER_BOB://kafka0:29092,LISTENER_FRED://localhost:9092
KAFKA_LISTENER_SECURITY_PROTOCOL_MAP: LISTENER_BOB:PLAINTEXT,LISTENER_FRED:PLAINTEXT
KAFKA_INTER_BROKER_LISTENER_NAME: LISTENER_BOB
```

tl;dr : You need to set advertised.listeners (or KAFKA\_ADVERTISED\_LISTENERS if you're using Docker images) to the external address (host/IP) so that clients can correctly connect to it. Otherwise they'll try to connect to the internal host address—and if that's not reachable then problems ensue.



You need to tell Kafka how the brokers can reach each other, but also make sure that external clients (producers/consumers) can reach the broker they need to.

The key thing is that when you run a client, the broker you pass to it is just where it's going to go and get the metadata about brokers in the cluster from. The actual host & IP that it will connect to for reading/writing data is based on **the data that the broker passes back in that initial connection**—even if it's just a single node and the broker returned is the same as the one connected to.

- KAFKA\_LISTENERS is a comma-separated list of listeners, and the host/ip and port to which Kafka binds to on which to listen. For more complex networking this might be an IP address associated with a given network interface on a machine. The default is 0.0.0.0, which means listening on all interfaces. (listeners)
- KAFKA\_ADVERTISED\_LISTENERS is a comma-separated list of listeners with their host/ip and port. This is the metadata that's passed back to clients. (advertised.listeners)
- KAFKA\_LISTENER\_SECURITY\_PROTOCOL\_MAP defines key/value pairs for the security protocol to use, per listener name. (listener.security.protocol.map)
- KAFKA\_INTER\_BROKER\_LISTENER\_NAME: Kafka brokers communicate between themselves, usually on the internal network (e.g. Docker network, AWS VPC, etc). To define which listener to use, specify KAFKA\_INTER\_BROKER\_LISTENER\_NAME (inter.broker.listener.name). The host/IP used must be accessible from the broker machine to others.